


## **VEBA:**

a modular end-to-end suite for *in silico* recovery, clustering, and analysis of prokaryotic, microeukaryotic, and viral genomes from metagenomes

**Josh L. Espinoza**, Ph.D  
Staff Scientist  
Dupont Lab

Department of Environment and Sustainability, *J. Craig Venter Institute*  
Department of Human Biology and Genomic Medicine, *J. Craig Venter Institute*



## Peer-Reviewed Publication



Research | [Open Access](#) | [Published: 12 October 2022](#)

# VEBA: a modular end-to-end suite for in silico recovery, clustering, and analysis of prokaryotic, microeukaryotic, and viral genomes from metagenomes


[Josh L. Espinoza](#)  & [Chris L. Dupont](#)

*BMC Bioinformatics* **23**, Article number: 419 (2022) | [Cite this article](#)

**85** Accesses | **8** Altmetric | [Metrics](#)

# Microbial ecology in the larger context

- Microorganisms provide insight into ecosystem resilience, sustainability, and human health
- Cataloguing and preserving biodiversity is paramount for discovering potential solutions to challenges we face as a growing civilization
- Metagenomics pertains to the *in silico* study of microorganisms within an ecosystem *in situ*
  - Most metagenomics suites have conflicting dependencies
  - Most metagenomics suites only support prokaryotes
    - *Candidate phyla radiation* (CPR) support require manual *post hoc* workflows
    - Limited support for viruses (if any)
    - Limited support for eukaryotes (if any)








## Why use *VEBA*?

- Directly recovers, quality assess, and classify prokaryotic, eukaryotic, and viral genomes from metagenomes/metatranscriptomes
- Automated handling of CPR
- Modular to accommodate multiple workflows
- Automates complex tasks in a user-friendly way
- Maximizes information gain from available data
- Can be used with co-assemblies or sample-specific assemblies (+ pseudo-coassemblies)
- Implements clustering at the species and protein level to make sample-specific genomes comparable across multiple samples
- All packages and dependencies are open-sourced (no complicated licensing)
- *VEBA* is installed in one command (i.e., `bash install_veba.sh`)
- *VEBA* database is downloaded/configured in one command (i.e., `bash download_databases.sh /path/to/veba_database`)

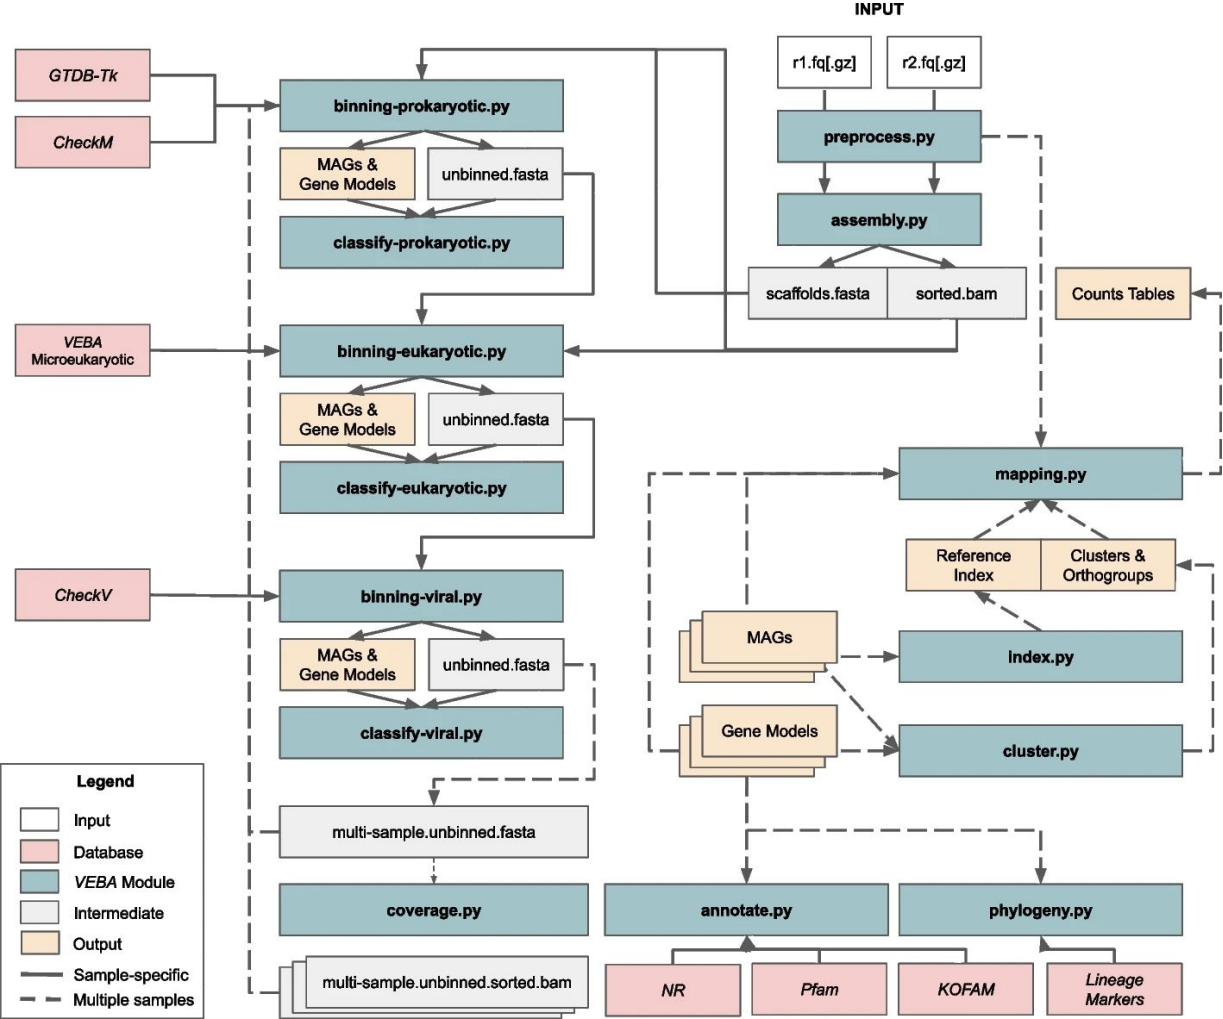
## VEBA Modules

- ***preprocess*** – Fastq quality trimming, adapter removal, decontamination, and read statistics calculations
- ***assembly*** – Assemble reads, align reads to assembly, and count mapped reads
- ***coverage*** – Align reads to (concatenated) reference and counts mapped reads
- ***binning-prokaryotic*** – Iterative consensus binning for recovering prokaryotic genomes with lineage-specific quality assessment
- ***binning-eukaryotic*** – Binning for recovering eukaryotic genomes with exon-aware gene modeling and lineage-specific quality assessment
- ***binning-viral*** – Detection of viral genomes and quality assessment
- ***classify-prokaryotic*** – Taxonomic classification and candidate phyla radiation adjusted quality
- ***classify-eukaryotic*** – Taxonomic classification of eukaryotic genomes
- ***classify-viral*** – Taxonomic classification and isolation source of viral genomes
- ***annotate*** – Annotates translated gene calls against NR, Pfam, and KOFAM
- ***cluster*** – Species-level clustering of genomes and lineage-specific orthogroup detection
- ***phylogeny*** – Constructs phylogenetic trees given a marker set
- ***index*** – Builds local or global index for alignment to genomes
- ***mapping*** – Aligns reads to local or global index of genomes

### Legend

	Preprocessing
	Identification
	Annotation
	Structural
	Quantification


# VEBA Modules



# What's the difference between co-assembly and sample-specific metagenomics?

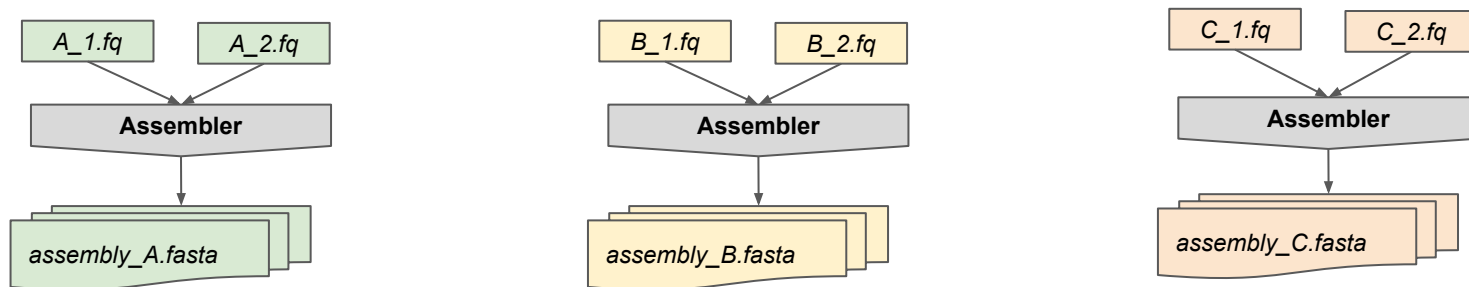
- **Co-assembly**

- Performing assembly when reads are from multiple samples (i.e., concatenated)



- **Sample-specific**

- Performing assembly for  $N$  samples individually resulting in  $N$  separate assemblies



# What's the difference between co-assembly and sample-specific metagenomics?

- **Co-assembly**

- **Pros:**

- Can increase read depth for low depth samples
    - Allows for direct comparison of features all samples
    - Coverage from multiple samples helps binning

- **Cons:**

- Can result in composite genomes that are not biologically accurate
    - Can result in many unbinned contigs
    - Requires much more computational power

- **Sample-specific assemblies**

- **Pros:**

- Can recover sample-specific individual strains instead of composites of multiple strains
    - Uses much less compute resources

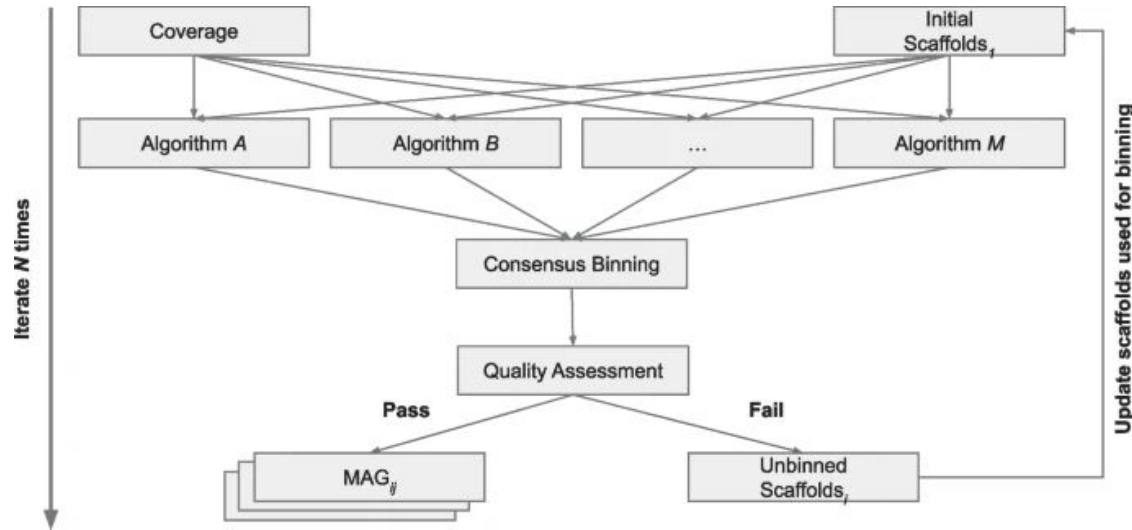
- **Cons:**

- Not as useful in low depth samples
    - Cannot directly compare abundances of biological features



## How does *VEBA* maximize available information for genome recovery?

- Uses iterative binning to feed unbinned contigs back into binning (currently, only implemented for prokaryotes)
- [Optional] Uses a “pseudo-coassembly” that makes use of unbinned contigs from multiple samples



- **Standardized parseable naming scheme:**
  - [SampleID]\_[Algorithm]\_DomainPrefix.[Iteration]\_[Name]
  - SRR17458623\_METABAT2\_P.1\_bin.1

- **P** for prokaryotes
- **E** for eukaryotes

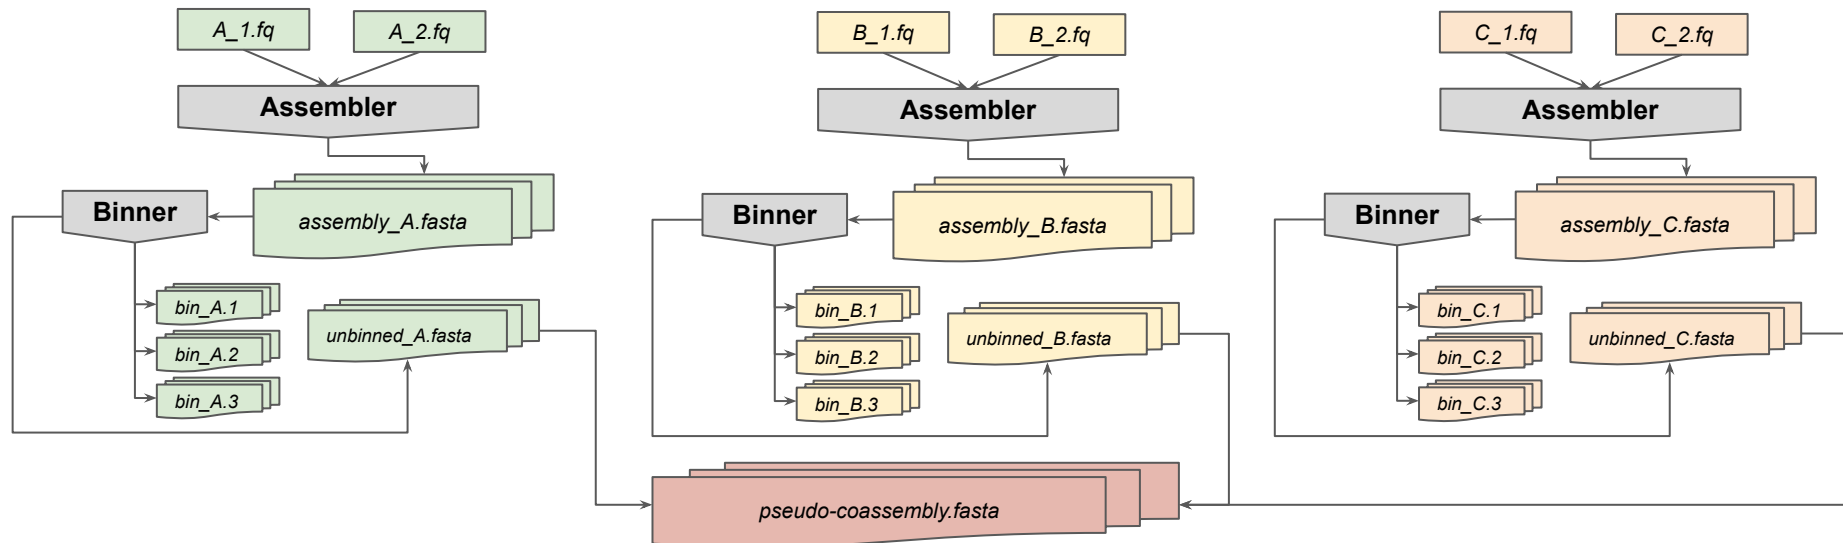
# What's the difference between *bona fide* co-assembly and pseudo-coassembly?

- **Co-assembly**

- Concatenate all forward reads (i.e., `cat *_1.fastq > concat_1.fastq`)
- Concatenate all reverse reads (i.e., `cat *_2.fastq > concat_2.fastq`)
- Assemble concatenated reads (i.e., `coassembly.fasta`)

- **Pseudo-coassembly**


- No additional assembly involved after sample-specific assembly
- Concatenate all unbinned contigs from sample-specific binning (i.e., `cat unbinned_*.fasta > pseudo-coassembly.fasta`)



# How does *VEBA* recover high-quality genomes from all domains?

- **Prokaryotes**

- **Recovery** - Iterative consensus binning (*MaxBin2*|*MetaBAT2*|*CONCOCT* → *DAS Tool*)
- **Gene calls** - *Prodigal* in metagenomics mode
- **Quality assess** - *CheckM* with automated workflow to handle CPR
- **Classification** - *GTDBTk*




- **Eukaryotes**

- **Recovery** - Binning using either *MetaBAT2* or *CONCOCT*
- **Gene calls** - *MetaEuk* exon-aware gene modeling using custom *VEBA* database
- **Validation** - *Tiara* to predict if genome is eukaryotic
- **Quality assess** - *BUSCO*
- **Classification** - *VEBA* sub-module that uses *MetaEuk* gene targets and bitscores



- **Viruses**

- **Recovery** - Binning using *VirFinder*
- **Gene calls** - *Prodigal* in metagenomics mode
- **Quality assess** - *CheckV*
- **Classification** - *VEBA* sub-module that uses *CheckV* references

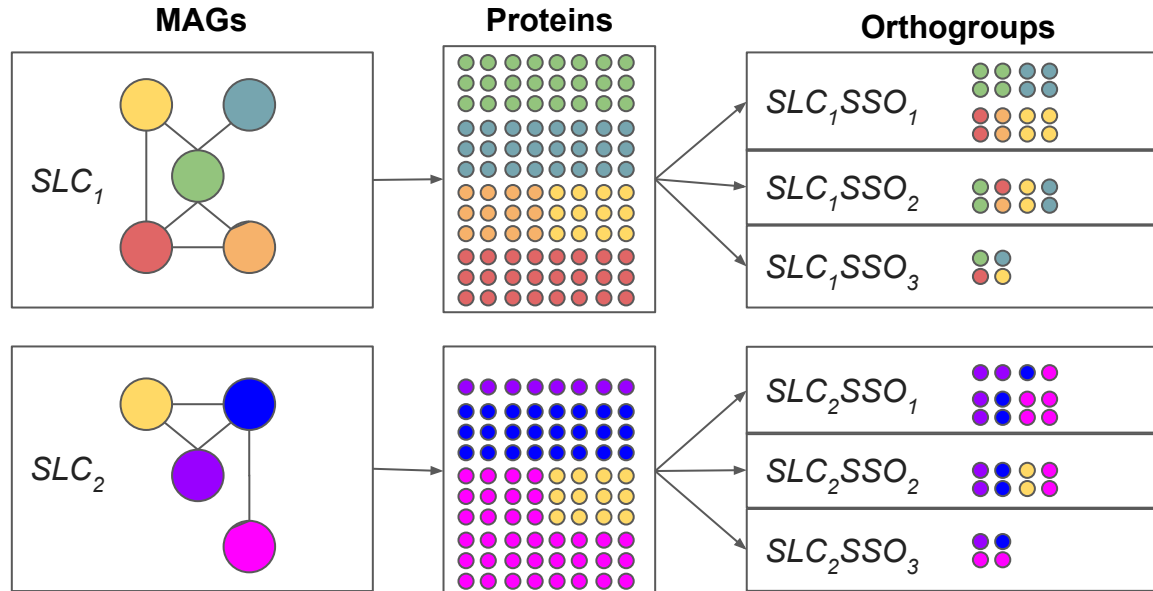


## Why does *VEBA* introduce YET ANOTHER eukaryotic database?!?

- **Current eukaryotic databases either do not target microeukaryotes and/or biased towards marine research**
  - ~~Useful for human microbiomes~~
  - ~~Useful for build microbiomes~~
- **Consensus database using microeukaryotic proteins from the following sources:**
  - *MMETSP*
  - *EukZoo*
  - *EukProt*
  - NCBI non-redundant
- **Used for gene modeling and taxonomy classification**
- **Streamlined by removing prokaryotic and higher eukaryotes**
- **Contains 48,006,918 proteins sequences from 42,922 unique species**
- **Available on FigShare (10.21 GB)**
  - [https://figshare.com/articles/dataset/Microeukaryotic\\_Protein\\_Database/19668855/1](https://figshare.com/articles/dataset/Microeukaryotic_Protein_Database/19668855/1)


# Clustering at the taxonomic and functional level

- **Taxonomic**
  - *Metagenome-assemble genomes* (MAG) are clustered by *average nucleotide identity* (ANI)
  - MAGs that cluster at 95% ANI are a *Species-level cluster* (SLC)
- **Functional**
  - All proteins within a SLC are clustered into *SLC-specific orthogroups* (SSO)





## Curse(s) of dimensionality




# Curse(s) of dimensionality

- The curse of dimensionality refers to various phenomena that arise when analyzing and organizing data in high-dimensional spaces
- *"As the number of features or dimensions grows, the amount of data we need to generalize accurately grows exponentially."*

- Charles Isbell, Professor and Senior Associate Dean, School of Interactive Computing, Georgia Tech

- Most NGS datasets (Number of features  $\gg$  number of observations)
  - Large P, Small N ( $P \gg N$ )
  - Machine learning algorithms tend to overfit
  - Relationships between observations can be misleading





# Using clustering for dimensionality reduction

- MAGs in the same cluster share highly similar genetic segments
- Mapping reads to MAGs in the same cluster will randomly assign read to one MAG
- Summing MAG counts w.r.t to SLC negates this randomness
- Same for *Genes* → SSOs

MAGs

Samples


	id_sample	SRR17458630	SRR17458614	SRR17458615	SRR17458638
	SRR17458614__METABAT2__E.1__bin.2	17	1209979	1168603	9765
	SRR17458615__METABAT2__E.1__bin.2	2	1041303	1346154	9613
	SRR17458630__METABAT2__E.1__bin.3	1474634	57	77	4
	SRR17458638__METABAT2__E.1__bin.2	19	408	450	821748
	SRR17458638__METABAT2__E.1__bin.3	33	4304	4556	4527193

Group MAGs by SLCs and sum counts

Samples


SLCs

	id_sample	SRR17458630	SRR17458614	SRR17458615	SRR17458638
	ESLC1	19	2251282	2514757	19378
	ESLC2	1474634	57	77	4
	ESLC3	19	408	450	821748
	ESLC4	33	4304	4556	4527193



## How can we put this all together?

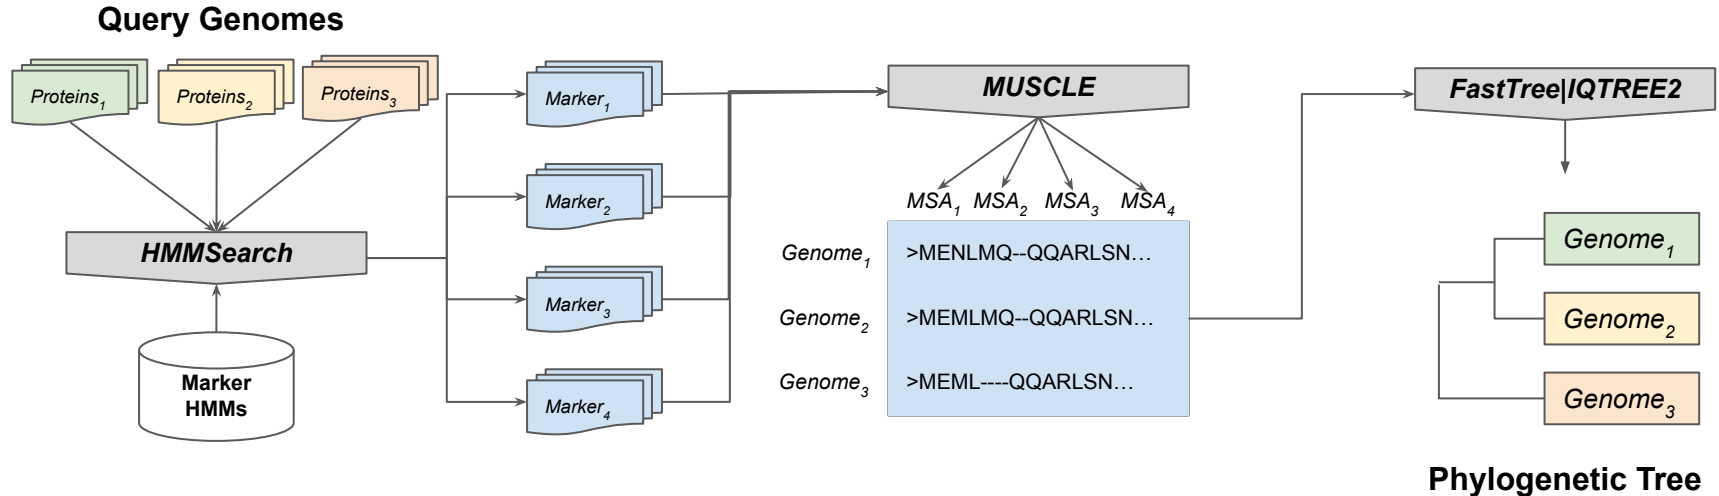
- Abundances/Expression comparable across samples
- Grouping features allows for high-level relationships between classes
- Multi-domain differential network analysis




# Phylogenetic inference using concatenated alignments

**VEBA includes the following HMM marker set:**

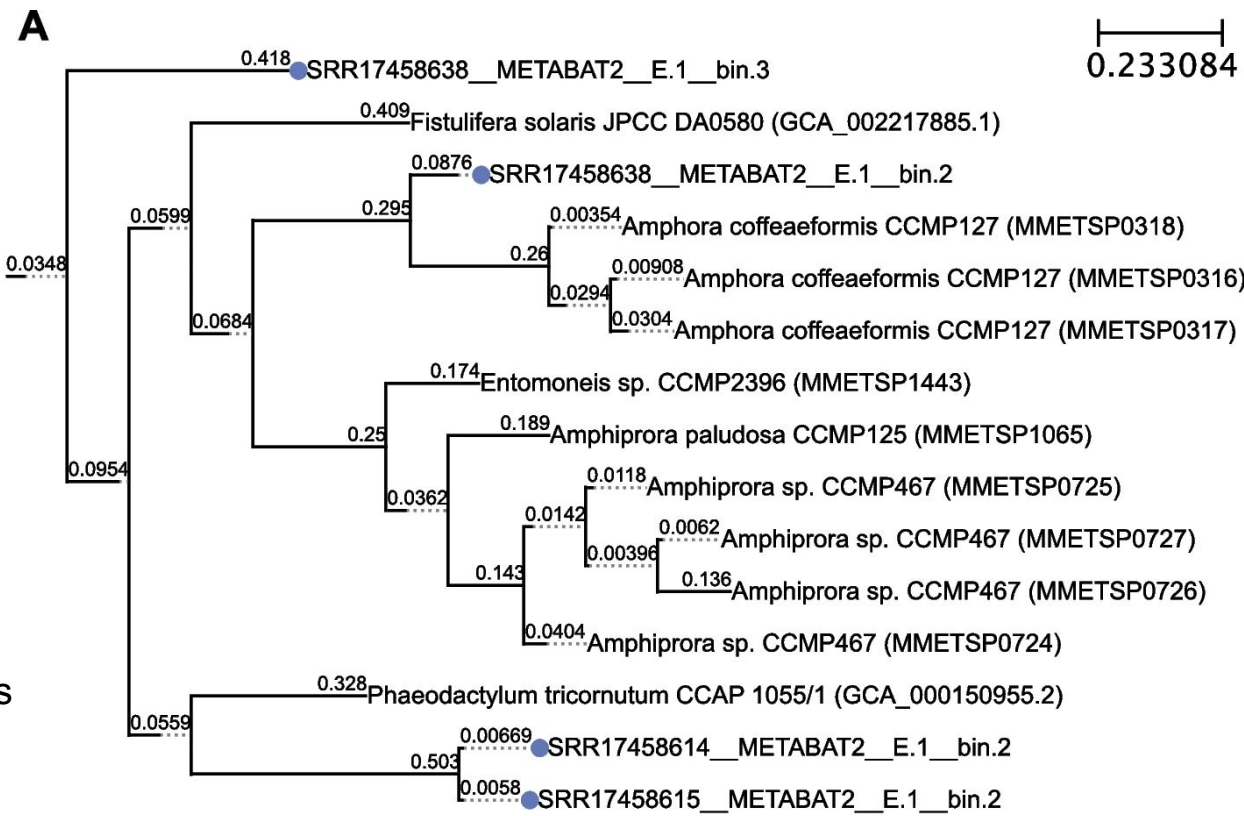
- Archaea\_76.hmm - (*Anvi'o*), Lee
- Bacteria\_71.hmm - (*Anvi'o*), Lee
- Protista\_83.hmm - (*Anvi'o*), Delmont
- Fungi\_593.hmm - (*FGMP*)
- CPR\_43.hmm - (*CheckM*)
- eukaryota\_odb10 - (*BUSCO*)



# Phylogenetic inference (IRL)



Taxonomic classifications reflect trends seen in phylogenetic inference



B

Diatom MAG	VEBA Eukaryotic Classification
SRR17458638__METABAT2__E.1__bin.3	c. Bacillariophyceae; o. Bacillariales; f. Bacillariaceae; g. s.
SRR17458614__METABAT2__E.1__bin.2	c. Bacillariophyceae; o. Naviculales; f. Phaeodactylaceae; g. Phaeodactylum; s. Phaeodactylum tricornutum [CCAP 1055/1]
SRR17458615__METABAT2__E.1__bin.2	c. Bacillariophyceae; o. Naviculales; f. Phaeodactylaceae; g. Phaeodactylum; s. Phaeodactylum tricornutum [CCAP 1055/1]
SRR17458638__METABAT2__E.1__bin.2	c. Bacillariophyceae; o. Thalassiophysales; f. Catenulaceae; g. Amphora; s.

## Tutorial walkthroughs for different meta-omics workflows

- **Downloading and preprocessing fastq files** - Explains how to download reads from NCBI and run *VEBA*'s `preprocess.py` module to decontaminate either metagenomic and/or metatranscriptomic reads.
- **Complete end-to-end metagenomics analysis** - Goes through assembling metagenomic reads, binning, clustering, classification, and annotation. We also show how to use the unbinned contigs in a pseudo-coassembly with guidelines on when it's a good idea to go this route.
- **Recovering viruses from metatranscriptomics** - Goes through assembling metatranscriptomic reads, viral binning, clustering, and classification.
- **Read mapping and counts tables** - Read mapping and generating counts tables at the contig, MAG, SLC, ORF, and SSO levels.
- **Phylogenetic inference** - Phylogenetic inference of eukaryotic diatoms.
- **Setting up *bona fide* co-assemblies for metagenomics or metatranscriptomics** - In the case where all samples are of low depth, it may be useful to use coassembly instead of sample-specific approaches. This walkthrough goes through concatenating reads, creating a reads table, coassembly of concatenated reads, aligning sample-specific reads to the coassembly for multiple sorted BAM files, and mapping reads for scaffold/transcript-level counts.

# Conclusion

- *VEBA* is a user-friendly metagenomics/metatranscriptomics software suite
- *VEBA* is all open-source so you don't have to deal with restrictive licenses (e.g., *GeneMark-EP+*)
- *VEBA* can handle:
  - Prokaryotes with direct support for CPR
  - Eukaryotes
  - Viruses
  - Sample-specific assemblies
  - Co-assemblies
  - Genomes from other pipelines/references
- *VEBA* is modular and can be used at many different stages of analysis
- *VEBA* has walkthroughs as step-by-step guides for different common workflows
- *VEBA* makes very complicated workflows extremely easy



## Questions?

- **E-mail**
  - [jespinoz@jcv.org](mailto:jespinoz@jcv.org)
- **LinkedIn:**
  - <https://www.linkedin.com/in/jolespin/>
- **Soothsayer Ecosystem (GitHub)**
  - **VEBA** (<https://github.com/jolespin/veba>)
    - A modular end-to-end suite for in silico recovery, clustering, and analysis of prokaryotic, microeukaryotic, and viral genomes from metagenomes
  - **Soothsayer** (<https://github.com/jolespin/soothsayer>)
    - High-level analysis package for (bio-)informatics
  - **Ensemble NetworkX** ([https://github.com/jolespin/ensemble\\_networkx](https://github.com/jolespin/ensemble_networkx))
    - Ensemble networks in Python
  - **Hive NetworkX** ([https://github.com/jolespin/hive\\_networkx](https://github.com/jolespin/hive_networkx))
    - Hive plots in Python
  - **Compositional** (<https://github.com/jolespin/compositional>)
    - Compositional data analysis in Python
  - **GenoPype** (<https://github.com/jolespin/genopype>)
    - Architecture for creating bash pipelines, in particular, for bioinformatics

## ***VEBA* Module Specifics**



## ***preprocess.py*** — Fastq quality trimming, adapter removal, and decontamination

- **Workflow:**
  - Wrapper around [\*fastq\\_preprocessor\*](#) (A “modernized” reimplementation of *KneadData*)
  - Automatic quality trimming and adapter removal and with *FastP*
  - [Optional] Removal/quantification of contamination based if reference provided:
    - *Bowtie2* - alignment based (e.g., removing human reads)
    - *BBDuk* - *k*-mer based (e.g., removing ribosomal reads)
    - Can quantify but not store read subsets (e.g., count ribosomal hits but don’t save them)
  - Calculate read statistics used *SeqKit*
- **Input:**
  - Raw paired reads (fastq)
- **Output:**
  - Verified quality trimmed reads (with contamination removed if applicable)
  - Summary statistics for full accounting of reads

## ***assembly.py*** — Assemble reads, align reads to assembly, and count mapped reads

- **Workflow:**
  - Assembles paired reads using **SPAdes**-based assemblers (e.g., ***metaSPAdes***, ***rnaSPAdes***)
  - Builds **Bowtie2** index and maps reads to assembly to produce sorted BAM file
  - Indexes sorted BAM file
  - Counts reads using ***featureCounts***
  - Calculates summary statistics with ***SeqKit***
- **Input:**
  - Raw paired reads (fastq)
- **Output:**
  - Assembly fasta (and **Bowtie2** index)
  - Sorted BAM (and ***Samtools*** index)
  - Summary statistics for assemblies (e.g., total bases, total contigs, N50, etc.)
  - Simplified Annotation Format [SAF] file used for ***featureCounts*** read counting

## ***coverage.py*** – Align reads to a reference and count mapped reads

- **Workflow:**
  - Aligned reads from different samples to a reference using **Bowtie2**
  - Produces multiple sorted BAM files and indexes
  - Counts reads using **featureCounts**
  - Calculate read statistics used **SeqKit**
  - [Optional] Only necessary if doing pseudo-coassembly
- **Input:**
  - Reference fasta
  - A table of read paths [id\_sample]<tab>[path/to/r1.fastq.gz]<tab>[path/to/r2.fastq.gz]
- **Output:**
  - Multiple sorted BAM (and **Samtools** indexes)
  - Summary statistics for full accounting of reads

## ***binning-prokaryotic.py*** — Iterative consensus binning for prokaryotes

- **Workflow:**
  - Calculated coverage tables needed for binning algorithms using **CoverM**
  - Models genes using **Prodigal**
  - Iterative binning:
    - A,B) **MaxBin2** (marker set 40,107); C) **MetaBAT2**; D) **CONCOCT**
    - **DAS Tool** (A,B,C,D) → Candidate binned genomes
    - Remove eukaryotic genomes classified by **Tiara** and genome size filter
    - **CheckM** → High quality metagenome assembled genomes [MAG]
  - **GTDB-Tk** to classify taxonomy
  - Reevaluated candidate phyla radiation [CPR] using **CheckM** CPR marker set
  - Calculate genome statistics using **SeqKit**
- **Input:**
  - Assembly fasta (i.e., scaffolds.fasta from assembly module)
  - Sorted BAM file
- **Output:**
  - MAG assemblies, cds, protein, gene models
  - Identifier tables (ORF  $\longleftrightarrow$  Contig  $\longleftrightarrow$  MAG)
  - Summary tables (genome statistics, quality metrics, and classifications)
  - ORF-level counts tables
  - Binned/Unbinned lists (useful for grepping) and unbinned fasta file (used for next step)

## ***binning-eukaryotic.py*** — Binning for recovering eukaryotic genomes

- **Workflow:**
  - Calculated coverage tables needed for binning algorithms using **CoverM**
  - Bin genomes using either **MetaBAT2** or **CONCOCT** (can't use both yet)
  - Remove prokaryotic genomes classified by **Tiara** and genome size filter
  - Exon-aware gene modeling using **MetaEuk** for candidate eukaryotic genomes
  - Lineage-specific quality assessment using **BUSCO** (Remove low quality genomes)
  - Calculate genome statistics using **SeqKit**
- **Input:**
  - Assembly fasta (i.e., unbinned.fasta from prokaryotic binning module)
  - Sorted BAM file
- **Output:**
  - MAG assemblies, cds, protein, gene models
  - Identifier tables (ORF  $\longleftrightarrow$  Contig  $\longleftrightarrow$  MAG)
  - Summary tables (genome statistics, quality metrics, **MetaEuk** targets, and classifications)
  - ORF-level counts tables
  - Binned/Unbinned lists (useful for grepping) and unbinned fasta file (used for next step)

## ***binning-viral.py*** — Binning for recovering viral genomes

- **Workflow:**
  - Identify candidate viral genomes using **VirFinder** (**geNomad** coming soon...)
  - Model genes using **Prodigal**
  - Quality assessment using **CheckV** (Remove low quality genomes)
  - Calculate genome statistics using **SeqKit**
- **Input:**
  - Assembly fasta (i.e., unbinned.fasta from prokaryotic binning module)
- **Output:**
  - MAG assemblies, cds, protein, gene models
  - Identifier tables (ORF  $\longleftrightarrow$  Contig  $\longleftrightarrow$  MAG)
  - Summary tables (genome statistics, quality metrics, isolation source, and classifications)
  - ORF-level counts tables
  - Binned/Unbinned lists (useful for grepping) and unbinned fasta file (useful for pseudo-coassembly)

## ***cluster.py*** — Species-level clustering of genomes and proteins

- **Workflow:**
  - Cluster MAGs by Average Nucleotide Identity [ANI] using ***FastANI***
  - For each species-level cluster [SLC]:
    - Cluster proteins into lineage-specific orthogroups via ***OrthoFinder***
- **Input:**
  - Scaffolds to bins table
  - List of genome paths
  - List of protein paths
- **Output:**
  - MAG assemblies, cds, protein, gene models
  - Identifier tables (Contig  $\longleftrightarrow$  MAG  $\longleftrightarrow$  SLC) & (ORF  $\longleftrightarrow$  Orthogroup)

## ***classify-prokaryotic.py*** — Taxonomic classification of prokaryotes

- **Workflow:**
  - Compiles ***GTDB-Tk*** classification files
- **Input:**
  - Prokaryotic binning directory
  - [Optional] Prokaryotic SLC clustering
- **Output:**
  - Taxonomy classifications for each MAG
  - [Optional] Prokaryotic cluster classification



## ***classify-eukaryotic.py*** — Taxonomic classification of eukaryotes

- **Workflow:**
  - Gets eukaryotic markers using ***HMMER***
  - Gets ***MetaEuk*** targets of eukaryotic markers
  - Classifies eukaryotic taxonomy based on bitscores and lineage
- **Input:**
  - Eukaryotic binning directory
  - [Optional] Eukaryotic SLC clustering
- **Output:**
  - Taxonomy classifications for each MAG
  - Gene source lineage with bitscores for each marker gene used in classification
  - [Optional] Eukaryotic cluster classification

## ***classify-viral.py*** — Taxonomic classification of viruses

- **Workflow:**
  - Use ***CheckV*** output and database to classify viruses and isolation source
- **Input:**
  - Viral binning directory
  - [Optional] Viral SLC clustering
- **Output:**
  - Taxonomy classifications for each MAG
  - [Optional] Viral cluster classification
  - [Optional] Consensus isolation source

## ***annotate.py*** — Annotates translated gene calls against NR, Pfam, and KOFAM

- **Workflow:**
  - Align proteins to NCBI's non-redundant database via ***Diamond***
  - Search for ***Pfam*** protein domains using ***HMMER***
  - Search for KEGG orthology using **KOFAMSCAN**
  - [Optional] Identifier mapping [id\_orf]<tab>[id\_contig]
- **Input:**
  - Protein fasta file
- **Output:**
  - Annotation table
  - [Optional] Contig-level annotations based solely on NR

## ***phylogeny.py*** — Constructs phylogenetic trees given a marker set

- **Workflow:**
  - Identifies marker proteins using **HMMER** based on user-provided database
  - [Optional] Remove hits based on marker-score thresholds
  - Protein alignment for each marker identified via **MUSCLE**
  - Alignments are trimmed using **ClipKIT**
  - Concatenate alignments
  - Approximately-maximum likelihood phylogenetic inference via **FastTree2**
  - [Optional] Maximum likelihood phylogenetic inference via **IQTREE2** (Takes a long time)
- **Input:**
  - Table of protein fasta files
  - HMM Database
  - [Optional] Table of marker score cutoffs
- **Output:**
  - Newick formatted phylogenetic tree
  - Concatenated multiple sequence alignment
  - Alignment table ( $n$  genomes,  $m$  markers,  $ij$ =fasta alignment)

## ***index.py*** — Builds index for alignment to genomes

- **Workflow:**
  - Creates reference index for binned genomes via ***Bowtie2***
  - Merges gene models (GFF3)
- **Input:**
  - Reference fasta file[s]
  - Gene model GFF3 file[s]
- **Output:**
  - Concatenated reference fasta
  - Concatenated reference fasta ***Bowtie2*** index
  - Concatenated gene models

## ***mapping.py*** — Builds index for alignment to genomes

- **Workflow:**

- Maps reads to **Bowtie2** reference index
- Counts reads for contigs and ORFs
- [Optional] Aggregates reads for MAG and SLC level
- [Optional] Calculates spatial coverage for each MAG (i.e., ratio of bases covered in genome)

- **Input:**

- Paired reads
- Reference index directory (contains index, fasta, GFF3, and SAF)
- [Optional] ORF to orthogroup identifier table
- [Optional] Contigs to MAG identifier table
- [Optional] Contigs to SLC identifier table

- **Output:**

- Sorted BAM file
- Paired unmapped reads
- ORF-level counts table, contig-level counts table
- [Optional] MAG-level counts table, SLC-level counts table, Orthogroup-level counts table