


VEBA:

a modular end-to-end suite for *in silico* recovery, clustering, and analysis of prokaryotic, microeukaryotic, and viral genomes from metagenomes

Josh L. Espinoza, Ph.D
Staff Scientist
Dupont Lab

Department of Environment and Sustainability, *J. Craig Venter Institute*
Department of Human Biology and Genomic Medicine, *J. Craig Venter Institute*



Peer-Reviewed Publication



Research | [Open Access](#) | [Published: 12 October 2022](#)

VEBA: a modular end-to-end suite for in silico recovery, clustering, and analysis of prokaryotic, microeukaryotic, and viral genomes from metagenomes


[Josh L. Espinoza](#)  & [Chris L. Dupont](#)

BMC Bioinformatics **23**, Article number: 419 (2022) | [Cite this article](#)

85 Accesses | **8** Altmetric | [Metrics](#)

Microbial ecology in the larger context

- Microorganisms provide insight into ecosystem resilience, sustainability, and human health
- Cataloguing and preserving biodiversity is paramount for discovering potential solutions to challenges we face as a growing civilization
- Metagenomics pertains to the *in silico* study of microorganisms within an ecosystem *in situ*
 - Most metagenomics suites have conflicting dependencies
 - Most metagenomics suites only support prokaryotes
 - *Candidate phyla radiation* (CPR) support require manual *post hoc* workflows
 - Limited support for viruses (if any)
 - Limited support for eukaryotes (if any)








Why use *VEBA*?

- Directly recovers, quality assess, and classify prokaryotic, eukaryotic, and viral genomes from metagenomes/metatranscriptomes
- Automated handling of CPR
- Modular to accommodate multiple workflows
- Automates complex tasks in a user-friendly way
- Maximizes information gain from available data
- Can be used with co-assemblies or sample-specific assemblies (+ pseudo-coassemblies)
- Implements clustering at the species and protein level to make sample-specific genomes comparable across multiple samples
- All packages and dependencies are open-sourced (no complicated licensing)
- *VEBA* is installed in one command (i.e., `bash install_veba.sh`)
- *VEBA* database is downloaded/configured in one command (i.e., `bash download_databases.sh /path/to/veba_database`)

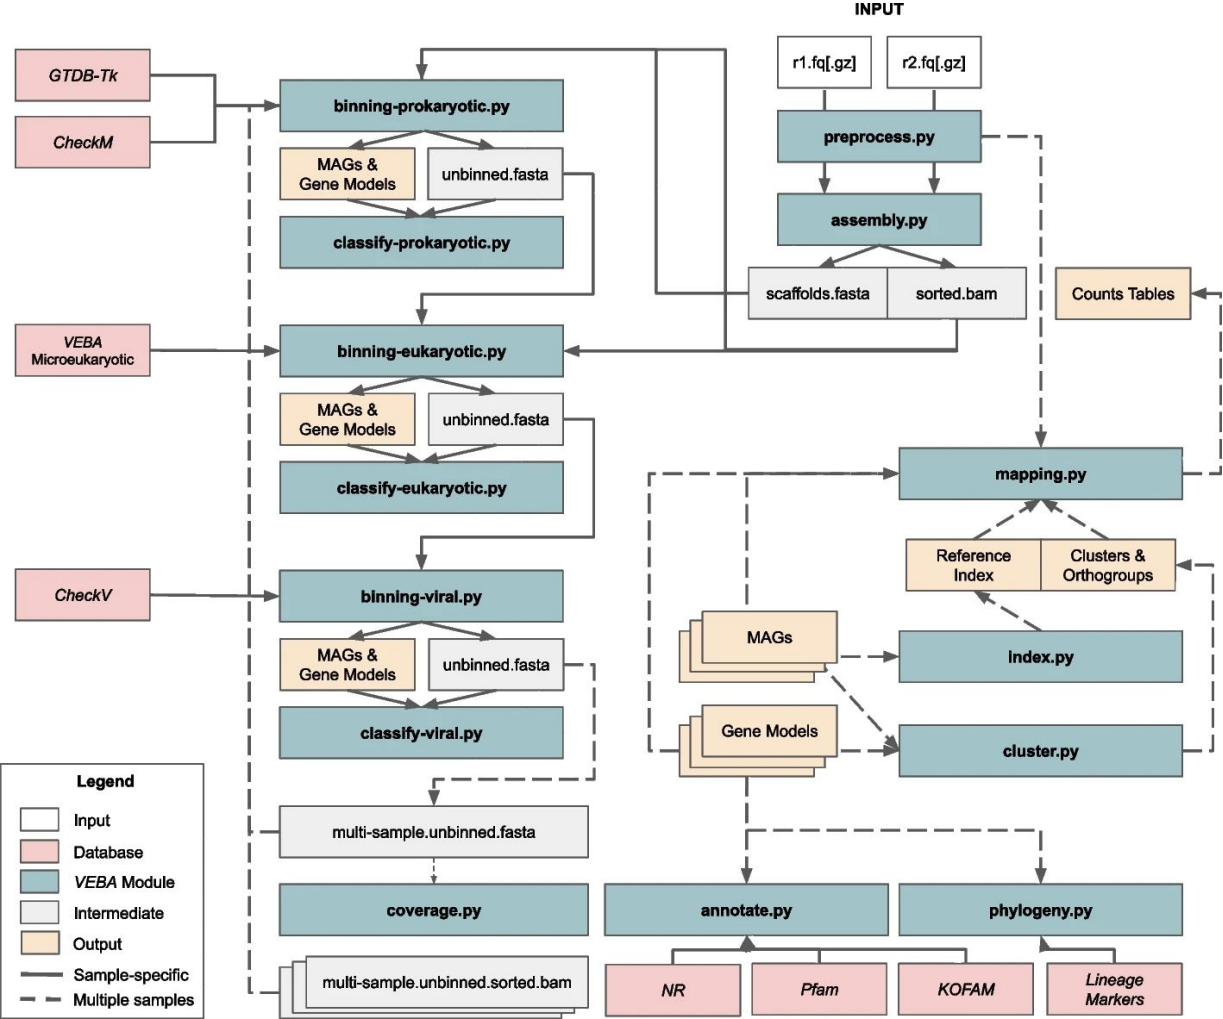
VEBA Modules

- ***preprocess*** – Fastq quality trimming, adapter removal, decontamination, and read statistics calculations
- ***assembly*** – Assemble reads, align reads to assembly, and count mapped reads
- ***coverage*** – Align reads to (concatenated) reference and counts mapped reads
- ***binning-prokaryotic*** – Iterative consensus binning for recovering prokaryotic genomes with lineage-specific quality assessment
- ***binning-eukaryotic*** – Binning for recovering eukaryotic genomes with exon-aware gene modeling and lineage-specific quality assessment
- ***binning-viral*** – Detection of viral genomes and quality assessment
- ***classify-prokaryotic*** – Taxonomic classification and candidate phyla radiation adjusted quality
- ***classify-eukaryotic*** – Taxonomic classification of eukaryotic genomes
- ***classify-viral*** – Taxonomic classification and isolation source of viral genomes
- ***annotate*** – Annotates translated gene calls against NR, Pfam, and KOFAM
- ***cluster*** – Species-level clustering of genomes and lineage-specific orthogroup detection
- ***phylogeny*** – Constructs phylogenetic trees given a marker set
- ***index*** – Builds local or global index for alignment to genomes
- ***mapping*** – Aligns reads to local or global index of genomes

Legend

	Preprocessing
	Identification
	Annotation
	Structural
	Quantification


VEBA Modules



What's the difference between co-assembly and sample-specific metagenomics?

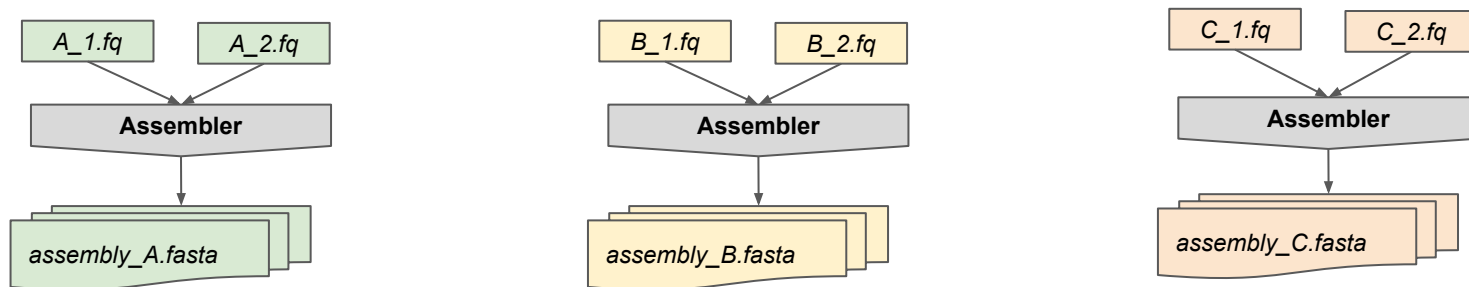
- **Co-assembly**

- Performing assembly when reads are from multiple samples (i.e., concatenated)



- **Sample-specific**

- Performing assembly for N samples individually resulting in N separate assemblies



What's the difference between co-assembly and sample-specific metagenomics?

- **Co-assembly**

- **Pros:**

- Can increase read depth for low depth samples
 - Allows for direct comparison of features all samples
 - Coverage from multiple samples helps binning

- **Cons:**

- Can result in composite genomes that are not biologically accurate
 - Can result in many unbinned contigs
 - Requires much more computational power

- **Sample-specific assemblies**

- **Pros:**

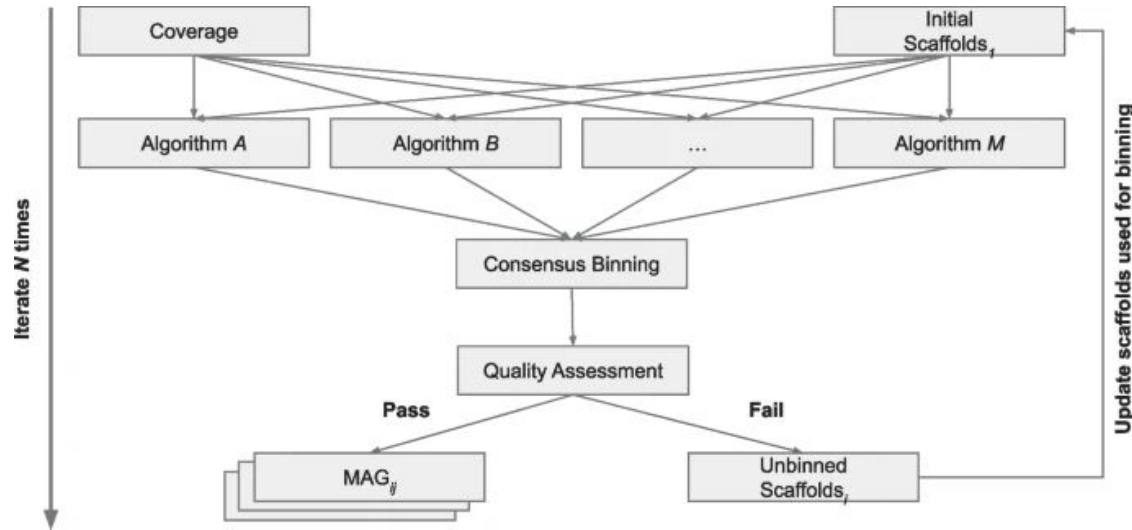
- Can recover sample-specific individual strains instead of composites of multiple strains
 - Uses much less compute resources

- **Cons:**

- Not as useful in low depth samples
 - Cannot directly compare abundances of biological features

How does *VEBA* maximize available information for genome recovery?

- Uses iterative binning to feed unbinned contigs back into binning (currently, only implemented for prokaryotes)
- [Optional] Uses a “pseudo-coassembly” that makes use of unbinned contigs from multiple samples



- **Standardized parseable naming scheme:**
 - [SampleID]_[Algorithm]_DomainPrefix.[Iteration]_[Name]
 - SRR17458623_METABAT2_P.1_bin.1

- **P** for prokaryotes
- **E** for eukaryotes

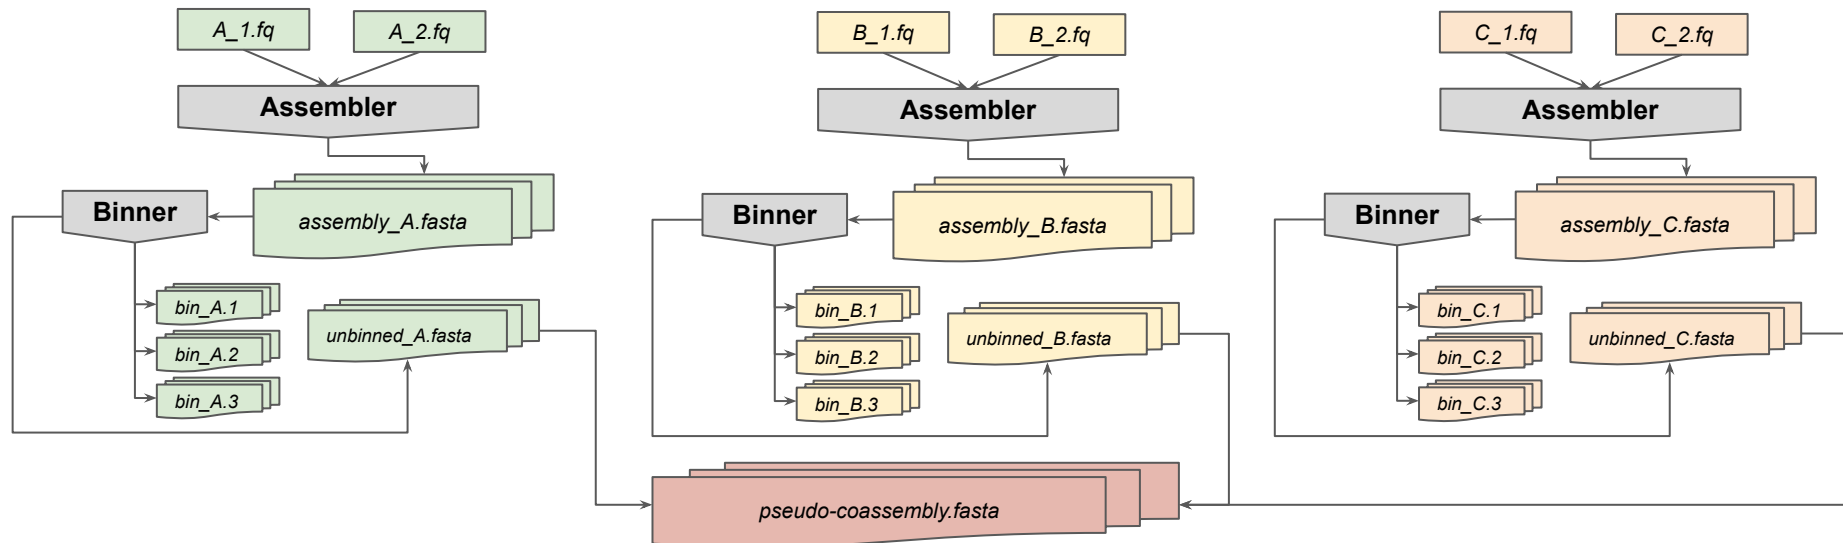
What's the difference between *bona fide* co-assembly and pseudo-coassembly?

- **Co-assembly**

- Concatenate all forward reads (i.e., `cat *_1.fastq > concat_1.fastq`)
- Concatenate all reverse reads (i.e., `cat *_2.fastq > concat_2.fastq`)
- Assemble concatenated reads (i.e., `coassembly.fasta`)

- **Pseudo-coassembly**


- No additional assembly involved after sample-specific assembly
- Concatenate all unbinned contigs from sample-specific binning (i.e., `cat unbinned_*.fasta > pseudo-coassembly.fasta`)



How does *VEBA* recover high-quality genomes from all domains?

- **Prokaryotes**

- **Recovery** - Iterative consensus binning (*MaxBin2*|*MetaBAT2*|*CONCOCT* → *DAS Tool*)
- **Gene calls** - *Prodigal* in metagenomics mode
- **Quality assess** - *CheckM* with automated workflow to handle CPR
- **Classification** - *GTDBTk*




- **Eukaryotes**

- **Recovery** - Binning using either *MetaBAT2* or *CONCOCT*
- **Gene calls** - *MetaEuk* exon-aware gene modeling using custom *VEBA* database
- **Validation** - *Tiara* to predict if genome is eukaryotic
- **Quality assess** - *BUSCO*
- **Classification** - *VEBA* sub-module that uses *MetaEuk* gene targets and bitscores



- **Viruses**

- **Recovery** - Binning using *VirFinder*
- **Gene calls** - *Prodigal* in metagenomics mode
- **Quality assess** - *CheckV*
- **Classification** - *VEBA* sub-module that uses *CheckV* references

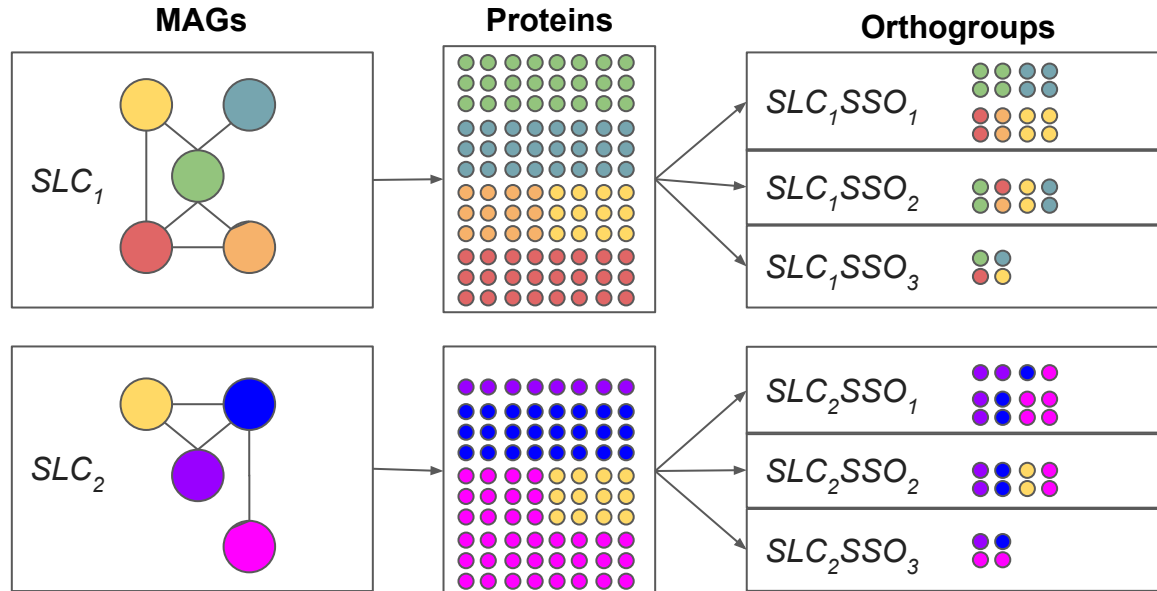


Why does *VEBA* introduce YET ANOTHER eukaryotic database?!?

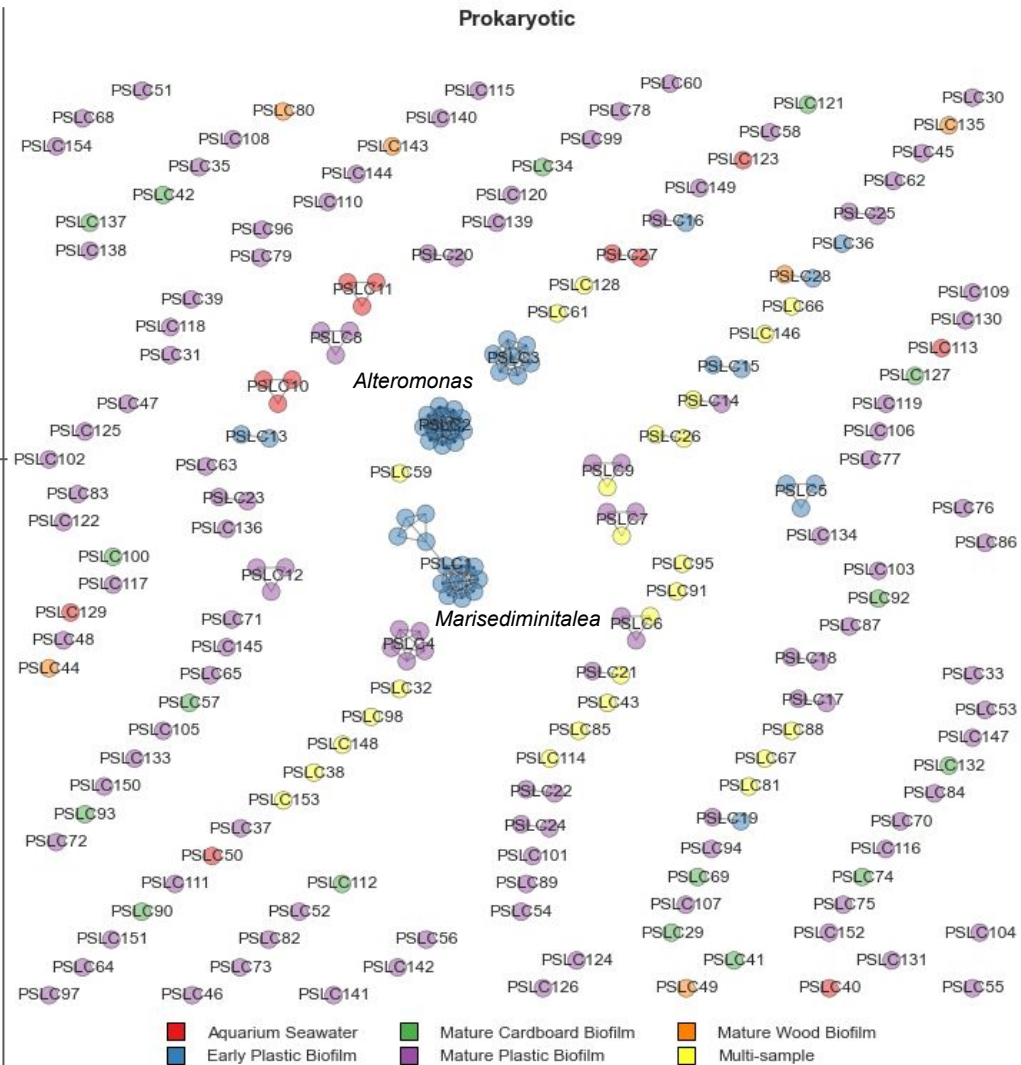


- **Current eukaryotic databases either do not target microeukaryotes and/or biased towards marine research**
 - ~~Useful for human microbiomes~~
 - ~~Useful for build microbiomes~~
- **Consensus database using microeukaryotic proteins from the following sources:**
 - *MMETSP*
 - *EukZoo*
 - *EukProt*
 - NCBI non-redundant
- **Used for gene modeling and taxonomy classification**
- **Streamlined by removing prokaryotic and higher eukaryotes**
- **Contains 48,006,918 proteins sequences from 42,922 unique species**
- **Available on FigShare (10.21 GB)**
 - https://figshare.com/articles/dataset/Microeukaryotic_Protein_Database/19668855/1

Clustering at the taxonomic and functional level


- **Taxonomic**
 - *Metagenome-assemble genomes* (MAG) are clustered by *average nucleotide identity* (ANI)
 - MAGs that cluster at 95% ANI are a *Species-level cluster* (SLC)
- **Functional**
 - All proteins within a SLC are clustered into *SLC-specific orthogroups* (SSO)



Species-level clustering (IRL)



Curse(s) of dimensionality




Curse(s) of dimensionality

- The curse of dimensionality refers to various phenomena that arise when analyzing and organizing data in high-dimensional spaces
- *"As the number of features or dimensions grows, the amount of data we need to generalize accurately grows exponentially."*

- Charles Isbell, Professor and Senior Associate Dean, School of Interactive Computing, Georgia Tech

- Most NGS datasets (Number of features \gg number of observations)
 - Large P, Small N ($P \gg N$)
 - Machine learning algorithms tend to overfit
 - Relationships between observations can be misleading



Using clustering for dimensionality reduction

- MAGs in the same cluster share highly similar genetic segments
- Mapping reads to MAGs in the same cluster will randomly assign read to one MAG
- Summing MAG counts w.r.t to SLC negates this randomness
- Same for *Genes* → SSOs

MAGs

Samples


	id_sample	SRR17458630	SRR17458614	SRR17458615	SRR17458638
	SRR17458614__METABAT2__E.1__bin.2	17	1209979	1168603	9765
	SRR17458615__METABAT2__E.1__bin.2	2	1041303	1346154	9613
	SRR17458630__METABAT2__E.1__bin.3	1474634	57	77	4
	SRR17458638__METABAT2__E.1__bin.2	19	408	450	821748
	SRR17458638__METABAT2__E.1__bin.3	33	4304	4556	4527193

Group MAGs by SLCs and sum counts

Samples


SLCs

	id_sample	SRR17458630	SRR17458614	SRR17458615	SRR17458638
	ESLC1	19	2251282	2514757	19378
	ESLC2	1474634	57	77	4
	ESLC3	19	408	450	821748
	ESLC4	33	4304	4556	4527193



How can we put this all together?

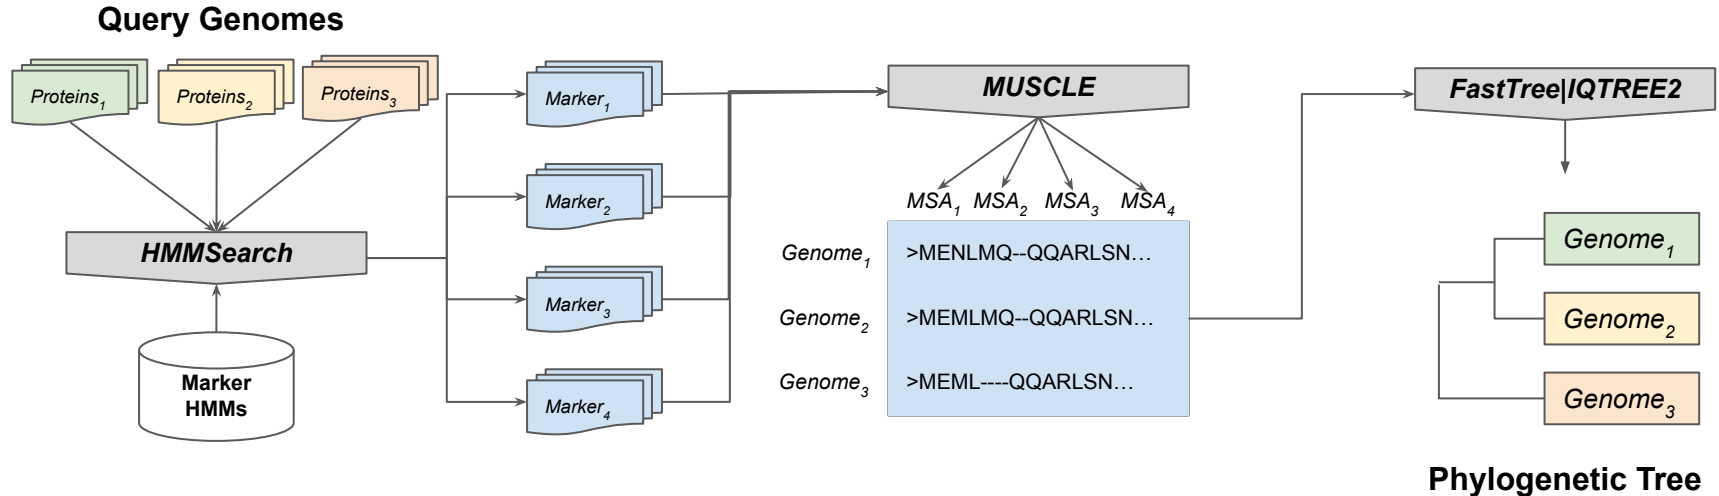
- Abundances/Expression comparable across samples
- Grouping features allows for high-level relationships between classes
- Multi-domain differential network analysis




Phylogenetic inference using concatenated alignments

VEBA includes the following HMM marker set:

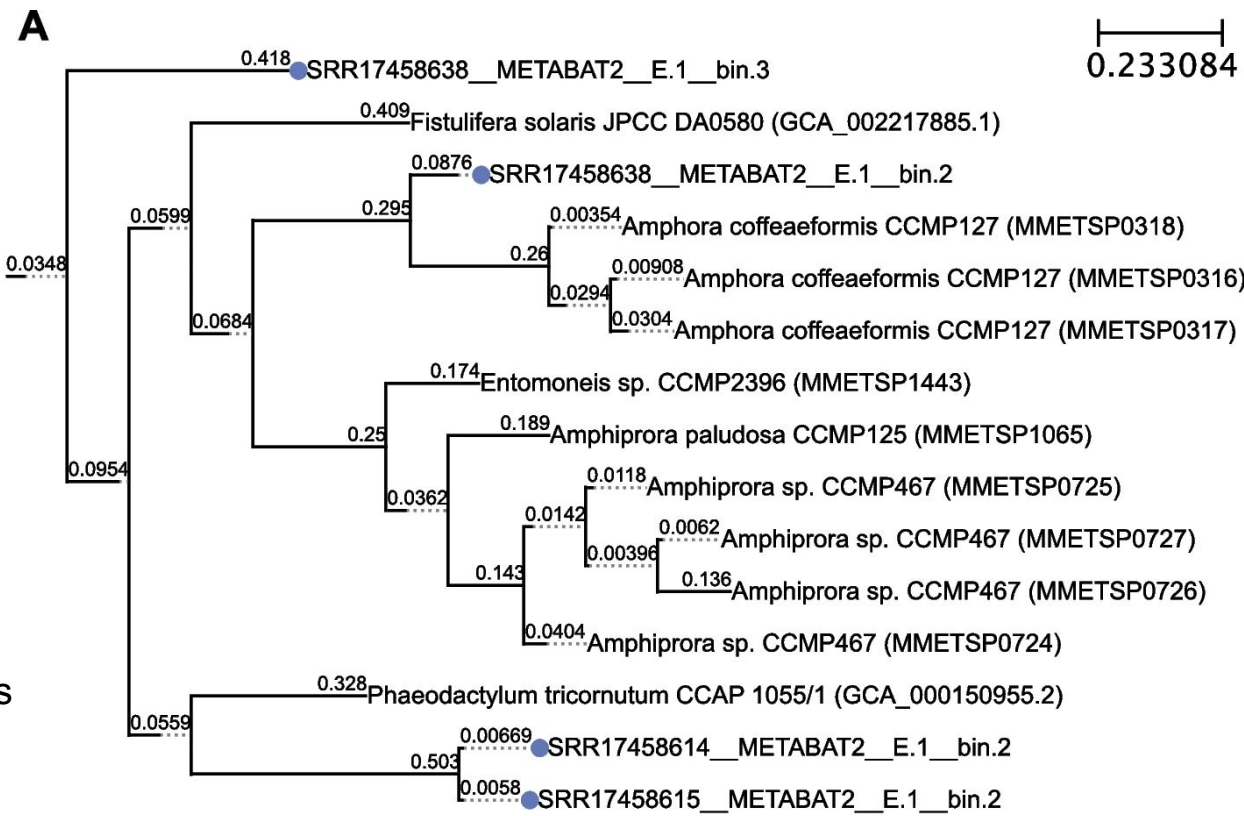
- Archaea_76.hmm - (*Anvi'o*), Lee
- Bacteria_71.hmm - (*Anvi'o*), Lee
- Protista_83.hmm - (*Anvi'o*), Delmont
- Fungi_593.hmm - (*FGMP*)
- CPR_43.hmm - (*CheckM*)
- eukaryota_odb10 - (*BUSCO*)



Phylogenetic inference (IRL)



Taxonomic classifications reflect trends seen in phylogenetic inference



B

Diatom MAG	VEBA Eukaryotic Classification
SRR17458638__METABAT2__E.1__bin.3	c. Bacillariophyceae;o. Bacillariales;f. Bacillariaceae;g. s
SRR17458614__METABAT2__E.1__bin.2	c. Bacillariophyceae;o. Naviculales;f. Phaeodactylaceae;g. Phaeodactylum;s. Phaeodactylum tricornutum [CCAP 1055/1]
SRR17458615__METABAT2__E.1__bin.2	c. Bacillariophyceae;o. Naviculales;f. Phaeodactylaceae;g. Phaeodactylum;s. Phaeodactylum tricornutum [CCAP 1055/1]
SRR17458638__METABAT2__E.1__bin.2	c. Bacillariophyceae;o. Thalassiophysales;f. Catenulaceae;g. Amphora;s

Tutorial walkthroughs for different meta-omics workflows

- **Downloading and preprocessing fastq files** - Explains how to download reads from NCBI and run *VEBA*'s `preprocess.py` module to decontaminate either metagenomic and/or metatranscriptomic reads.
- **Complete end-to-end metagenomics analysis** - Goes through assembling metagenomic reads, binning, clustering, classification, and annotation. We also show how to use the unbinned contigs in a pseudo-coassembly with guidelines on when it's a good idea to go this route.
- **Recovering viruses from metatranscriptomics** - Goes through assembling metatranscriptomic reads, viral binning, clustering, and classification.
- **Read mapping and counts tables** - Read mapping and generating counts tables at the contig, MAG, SLC, ORF, and SSO levels.
- **Phylogenetic inference** - Phylogenetic inference of eukaryotic diatoms.
- **Setting up *bona fide* co-assemblies for metagenomics or metatranscriptomics** - In the case where all samples are of low depth, it may be useful to use coassembly instead of sample-specific approaches. This walkthrough goes through concatenating reads, creating a reads table, coassembly of concatenated reads, aligning sample-specific reads to the coassembly for multiple sorted BAM files, and mapping reads for scaffold/transcript-level counts.

Conclusion

- *VEBA* is a user-friendly metagenomics/metatranscriptomics software suite
- *VEBA* is all open-source so you don't have to deal with annoying licenses (e.g., *GeneMark-EP+*)
- *VEBA* can handle:
 - Prokaryotes with direct support for CPR
 - Eukaryotes
 - Viruses
 - Sample-specific assemblies
 - Co-assemblies
 - Genomes from other pipelines/references
- *VEBA* is modular and can be used at many different stages of analysis
- *VEBA* has walkthroughs as step-by-step guides for different common workflows
- *VEBA* makes very complicated workflows extremely easy



Questions?

- **E-mail**
 - jespinoz@jcv.org
- **LinkedIn:**
 - <https://www.linkedin.com/in/jolespin/>
- **Soothsayer Ecosystem (GitHub)**
 - **VEBA** (<https://github.com/jolespin/veba>)
 - A modular end-to-end suite for in silico recovery, clustering, and analysis of prokaryotic, microeukaryotic, and viral genomes from metagenomes
 - **Soothsayer** (<https://github.com/jolespin/soothsayer>)
 - High-level analysis package for (bio-)informatics
 - **Ensemble NetworkX** (https://github.com/jolespin/ensemble_networkx)
 - Ensemble networks in Python
 - **Hive NetworkX** (https://github.com/jolespin/hive_networkx)
 - Hive plots in Python
 - **Compositional** (<https://github.com/jolespin/compositional>)
 - Compositional data analysis in Python
 - **GenoPype** (<https://github.com/jolespin/genopype>)
 - Architecture for creating bash pipelines, in particular, for bioinformatics

***VEBA* Module Specifics**

preprocess.py — Fastq quality trimming, adapter removal, and decontamination

- **Workflow:**

- Wrapper around [*fastq_preprocessor*](#) (A “modernized” reimplementation of *KneadData*)
- Automatic quality trimming and adapter removal and with *FastP*
- [Optional] Removal/quantification of contamination based if reference provided:
 - *Bowtie2* - alignment based (e.g., removing human reads)
 - *BBDuk* - *k*-mer based (e.g., removing ribosomal reads)
 - Can quantify but not store read subsets (e.g., count ribosomal hits but don’t save them)
- Calculate read statistics used *SeqKit*

- **Input:**

- Raw paired reads (fastq)

- **Output:**

- Verified quality trimmed reads (with contamination removed if applicable)
- Summary statistics for full accounting of reads

assembly.py — Assemble reads, align reads to assembly, and count mapped reads

- **Workflow:**
 - Assembles paired reads using **SPAdes**-based assemblers (e.g., **metaSPAdes**, **rnaSPAdes**)
 - Builds **Bowtie2** index and maps reads to assembly to produce sorted BAM file
 - Indexes sorted BAM file
 - Counts reads using **featureCounts**
 - Calculates summary statistics with **SeqKit**
- **Input:**
 - Raw paired reads (fastq)
- **Output:**
 - Assembly fasta (and **Bowtie2** index)
 - Sorted BAM (and **Samtools** index)
 - Summary statistics for assemblies (e.g., total bases, total contigs, N50, etc.)
 - Simplified Annotation Format [SAF] file used for **featureCounts** read counting

coverage.py — Align reads to a reference and count mapped reads

- **Workflow:**
 - Aligned reads from different samples to a reference using **Bowtie2**
 - Produces multiple sorted BAM files and indexes
 - Counts reads using **featureCounts**
 - Calculate read statistics used **SeqKit**
 - [Optional] Only necessary if doing pseudo-coassembly
- **Input:**
 - Reference fasta
 - A table of read paths [id_sample]<tab>[path/to/r1.fastq.gz]<tab>[path/to/r2.fastq.gz]
- **Output:**
 - Multiple sorted BAM (and **Samtools** indexes)
 - Summary statistics for full accounting of reads

binning-prokaryotic.py — Iterative consensus binning for prokaryotes

- **Workflow:**
 - Calculated coverage tables needed for binning algorithms using ***CoverM***
 - Models genes using ***Prodigal***
 - Iterative binning:
 - A,B) ***MaxBin2*** (marker set 40,107); C) ***MetaBAT2***; D) ***CONCOCT***
 - ***DAS Tool*** (A,B,C,D) → Candidate binned genomes
 - Remove eukaryotic genomes classified by ***Tiara*** and genome size filter
 - ***CheckM*** → High quality metagenome assembled genomes [MAG]
 - ***GTDB-Tk*** to classify taxonomy
 - Reevaluated candidate phyla radiation [CPR] using ***CheckM*** CPR marker set
 - Calculate genome statistics using ***SeqKit***
- **Input:**
 - Assembly fasta (i.e., scaffolds.fasta from assembly module)
 - Sorted BAM file
- **Output:**
 - MAG assemblies, cds, protein, gene models
 - Identifier tables (ORF \longleftrightarrow Contig \longleftrightarrow MAG)
 - Summary tables (genome statistics, quality metrics, and classifications)
 - ORF-level counts tables
 - Binned/Unbinned lists (useful for grepping) and unbinned fasta file (used for next step)

binning-eukaryotic.py — Binning for recovering eukaryotic genomes

- **Workflow:**
 - Calculated coverage tables needed for binning algorithms using **CoverM**
 - Bin genomes using either **MetaBAT2** or **CONCOCT** (can't use both yet)
 - Remove prokaryotic genomes classified by **Tiara** and genome size filter
 - Exon-aware gene modeling using **MetaEuk** for candidate eukaryotic genomes
 - Lineage-specific quality assessment using **BUSCO** (Remove low quality genomes)
 - Calculate genome statistics using **SeqKit**
- **Input:**
 - Assembly fasta (i.e., unbinned.fasta from prokaryotic binning module)
 - Sorted BAM file
- **Output:**
 - MAG assemblies, cds, protein, gene models
 - Identifier tables (ORF \longleftrightarrow Contig \longleftrightarrow MAG)
 - Summary tables (genome statistics, quality metrics, **MetaEuk** targets, and classifications)
 - ORF-level counts tables
 - Binned/Unbinned lists (useful for grepping) and unbinned fasta file (used for next step)

binning-viral.py — Binning for recovering viral genomes

- **Workflow:**
 - Identify candidate viral genomes using **VirFinder** (**geNomad** coming soon...)
 - Model genes using **Prodigal**
 - Quality assessment using **CheckV** (Remove low quality genomes)
 - Calculate genome statistics using **SeqKit**
- **Input:**
 - Assembly fasta (i.e., unbinned.fasta from prokaryotic binning module)
- **Output:**
 - MAG assemblies, cds, protein, gene models
 - Identifier tables (ORF \longleftrightarrow Contig \longleftrightarrow MAG)
 - Summary tables (genome statistics, quality metrics, isolation source, and classifications)
 - ORF-level counts tables
 - Binned/Unbinned lists (useful for grepping) and unbinned fasta file (useful for pseudo-coassembly)

cluster.py — Species-level clustering of genomes and proteins

- **Workflow:**
 - Cluster MAGs by Average Nucleotide Identity [ANI] using ***FastANI***
 - For each species-level cluster [SLC]:
 - Cluster proteins into lineage-specific orthogroups via ***OrthoFinder***
- **Input:**
 - Scaffolds to bins table
 - List of genome paths
 - List of protein paths
- **Output:**
 - MAG assemblies, cds, protein, gene models
 - Identifier tables (Contig \longleftrightarrow MAG \longleftrightarrow SLC) & (ORF \longleftrightarrow Orthogroup)

classify-prokaryotic.py — Taxonomic classification of prokaryotes

- **Workflow:**
 - Compiles ***GTDB-Tk*** classification files
- **Input:**
 - Prokaryotic binning directory
 - [Optional] Prokaryotic SLC clustering
- **Output:**
 - Taxonomy classifications for each MAG
 - [Optional] Prokaryotic cluster classification

classify-eukaryotic.py — Taxonomic classification of eukaryotes

- **Workflow:**
 - Gets eukaryotic markers using **HMMER**
 - Gets **MetaEuk** targets of eukaryotic markers
 - Classifies eukaryotic taxonomy based on bitscores and lineage
- **Input:**
 - Eukaryotic binning directory
 - [Optional] Eukaryotic SLC clustering
- **Output:**
 - Taxonomy classifications for each MAG
 - Gene source lineage with bitscores for each marker gene used in classification
 - [Optional] Eukaryotic cluster classification

classify-viral.py — Taxonomic classification of viruses

- **Workflow:**
 - Use ***CheckV*** output and database to classify viruses and isolation source
- **Input:**
 - Viral binning directory
 - [Optional] Viral SLC clustering
- **Output:**
 - Taxonomy classifications for each MAG
 - [Optional] Viral cluster classification
 - [Optional] Consensus isolation source

annotate.py — Annotates translated gene calls against NR, Pfam, and KOFAM

- **Workflow:**
 - Align proteins to NCBI's non-redundant database via ***Diamond***
 - Search for ***Pfam*** protein domains using ***HMMER***
 - Search for KEGG orthology using **KOFAMSCAN**
 - [Optional] Identifier mapping [id_orf]<tab>[id_contig]
- **Input:**
 - Protein fasta file
- **Output:**
 - Annotation table
 - [Optional] Contig-level annotations based solely on NR

phylogeny.py — Constructs phylogenetic trees given a marker set

- **Workflow:**
 - Identifies marker proteins using **HMMER** based on user-provided database
 - [Optional] Remove hits based on marker-score thresholds
 - Protein alignment for each marker identified via **MUSCLE**
 - Alignments are trimmed using **ClipKIT**
 - Concatenate alignments
 - Approximately-maximum likelihood phylogenetic inference via **FastTree2**
 - [Optional] Maximum likelihood phylogenetic inference via **IQTREE2** (Takes a long time)
- **Input:**
 - Table of protein fasta files
 - HMM Database
 - [Optional] Table of marker score cutoffs
- **Output:**
 - Newick formatted phylogenetic tree
 - Concatenated multiple sequence alignment
 - Alignment table (n genomes, m markers, ij =fasta alignment)

index.py — Builds index for alignment to genomes

- **Workflow:**
 - Creates reference index for binned genomes via ***Bowtie2***
 - Merges gene models (GFF3)
- **Input:**
 - Reference fasta file[s]
 - Gene model GFF3 file[s]
- **Output:**
 - Concatenated reference fasta
 - Concatenated reference fasta ***Bowtie2*** index
 - Concatenated gene models

mapping.py — Builds index for alignment to genomes

- **Workflow:**

- Maps reads to **Bowtie2** reference index
- Counts reads for contigs and ORFs
- [Optional] Aggregates reads for MAG and SLC level
- [Optional] Calculates spatial coverage for each MAG (i.e., ratio of bases covered in genome)

- **Input:**

- Paired reads
- Reference index directory (contains index, fasta, GFF3, and SAF)
- [Optional] ORF to orthogroup identifier table
- [Optional] Contigs to MAG identifier table
- [Optional] Contigs to SLC identifier table

- **Output:**

- Sorted BAM file
- Paired unmapped reads
- ORF-level counts table, contig-level counts table
- [Optional] MAG-level counts table, SLC-level counts table, Orthogroup-level counts table